

DevRev AI Agent 007: Tooling up for Success

Team - 41

Inter-IIT Final Report

Contents

1	Introduction	2
2	Approaches	2
2.1	Chain of Thought	2
2.2	ReAct	3
2.3	ToolLLM with ChatGPT backbone	5
2.3.1	API-Retrieval	5
2.3.2	Depth First Search-based Decision Tree (DFSDT)	5
2.4	Few-shot Prompting	7
3	Results and Inferences	8
3.1	ToolEval	8
3.1.1	Difference between CoT and ReAct	10
3.1.2	Analysis of Model Usage: Tokens and Time	10
3.2	Reasoning on the Issues with ToolLLM (ChatGPT Backbone)	10
4	New Tools (Bonus Part)	11
5	Future Work	13
6	Conclusion	13

1 Introduction

The intersection of Natural Language Processing (NLP) and Domain-Specific question-answering has gained a lot of attention with the emergence of Large Language Models (LLMs) and Large Multimodal Models (LMMs). This emerging area focuses on empowering intelligent agents to comprehend and respond to user queries in specific domains, pushing the boundaries of the typical human-computer interaction.

The competition revolves around the dynamic interplay between the Language Model (L), the set of tools (T), and the user queries (Q). In a conversational format, users engage with the AI agent, generating questions that span the spectrum of tasks within the product's lifecycle. From basic inquiries to nuanced requests, the AI Agent 007 is tasked with dynamically selecting and utilizing tools to provide comprehensive and meaningful responses. Notably, the set of tools T is dynamic; it is subject to evolution, with new tools introduced and existing ones undergoing modifications or removals. Consequently, the AI agent must exhibit graceful adaptability to these changes.

2 Approaches

2.1 Chain of Thought

Chain of thought prompting [1] is a sequential series of interconnected ideas or concepts that stimulate cognitive processes, leading to a continuous flow of reflective or analytical thinking. It is a series of intermediate reasoning steps. In this case the input prompt is given the information about tools, and 3 examples which includes Question, Sequence of tools to choose and Output in JSON format.

Imagine you're trying to solve a puzzle. You wouldn't jump straight to the answer without thinking, right? You'd go through a series of steps, analyzing different options and building your logic step-by-step. This is exactly what Chain of Thought Prompting (CoT) does for large language models (LLMs).

CoT is like a mental map that helps the LLM navigate complex problems. Instead of simply being told the answer, the LLM is shown a series of "intermediate reasoning steps" that guide it towards the solution. This process allows the LLM to develop its own understanding of the problem and why the answer makes sense.

Here's an analogy to understand it better:

Imagine you are cooking a recipe for the first time. You wouldn't just throw all the ingredients together and hope for the best, right? Instead, you would follow the recipe step-by-step, adding each ingredient in the right order and at the right time.

CoT works in the same way. It guides the LLM through the reasoning process by providing step-by-step instructions. This helps the LLM to develop a deeper understanding of the problem and how to solve it, just like you would with a recipe.

Here are some of the benefits of CoT:

1. **Improved reasoning:** CoT helps LLMs to tackle complex problems that they would not be able to solve otherwise. This includes tasks like arithmetic, commonsense reasoning, and symbolic reasoning.
2. **More transparent reasoning:** CoT allows us to see how the LLM is thinking and why it came to a particular conclusion. This is helpful for debugging and improving the LLM's performance.
3. **More flexible reasoning:** CoT allows LLMs to adapt their reasoning process to different situations. This makes them more versatile and able to handle a wider range of problems.

So, basically in this technique a series of intermediate reasoning steps—significantly improves the ability of large language models to perform complex reasoning. In particular, reasoning abilities emerge naturally in sufficiently large language models via a simple method called chain-of-thought prompting, where a few chain of thought demonstrations are provided as exemplars in prompting. Chain-of-thought prompting enables large language models to tackle complex arithmetic, commonsense, and symbolic reasoning tasks. Example of such prompting is :

Question: Prioritize my P0 issues and add them to the current sprint

Chain of thoughts:

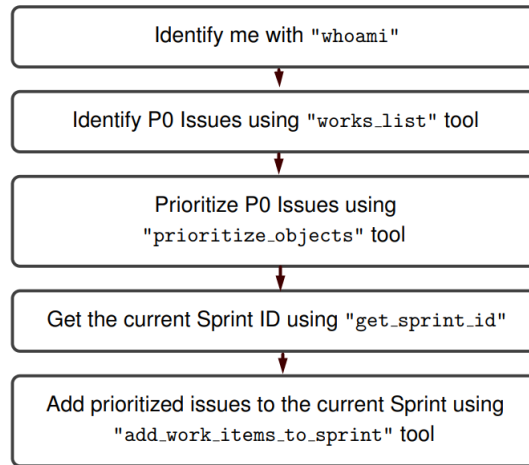


Figure 1: Chain of thoughts: Sequence of tools provided in prompt

An example illustrating the same is as follows:

Question : List all high severity tickets coming in from slack from customer Cust123 and generate a summary of them.

Output

<pre>[{ "tool_name": "search_object_by_name", "arguments": [{ "argument_name": "query", "argument_value": "Cust123" }] }, { "tool_name": "works_list", "arguments": [{ "argument_name": "ticket.rev_org", "argument_value": ["\$\$PREV[0]"] }, { "argument_name": "ticket.severity", "argument_value": ["high"] }] },]</pre>	<pre>{ "argument_name": "ticket.source_channel", "argument_value": ["slack"] }, { "argument_name": "type", "argument_value": ["ticket"] }], { "tool_name": "summarize_objects", "arguments": [{ "argument_name": "objects", "argument_value": "\$\$PREV[1]" }] }]</pre>
--	---

Figure 2: COT output which is exact same as provided in the Problem statement

2.2 ReAct

ReAct [2] is an interactive agent framework addressing complex policy learning challenges by expanding the action space to include language actions. Leveraging a substantial language model, it excels in reasoning tasks such as question answering and fact verification, demonstrating flexibility, generality, and robustness. The hybrid

strategy combines internal and external knowledge for optimal performance. As mentioned in the mid-term report, we generated data points using the **ReAct** framework [3], which will be tested using **ToolEval**.

Working:

- The framework was fed certain fewshot examples (path of file: in the following format (below) ToolQA+React/benchmarks/ReAct/code/fewshots.py).
- The tools and their descriptions were fed in a prompts file (path of file: ToolQA+React/benchmarks/ReAct/code/prompts.py).
- Input to the code is given in the file *qapairs.jsonl* (path of file: ToolQA+React/data/qapairs.jsonl).
- The file tests.py (path of file: ToolQA+React/benchmarks/ReAct/code/tests.py) was executed, where a **ReactAgent** was simulated and gpt-3.5-turbo model was prompted to give as output a set of thoughts, actions and observations as in the example below.

Query 1

Summarize high severity tickets from the customer UltimateCustomer.

Output:

Thought 1: This is a question related to summarizing high severity tickets from the customer UltimateCustomer.

Action 1: Response[{"tool_name": "search_object_by_name", "arguments": [{"argument_name": "query", "argument_value": "UltimateCustomer"}]}]

Observation 1: Called the function search_object_by_name and obtained the object of UltimateCustomer.

Thought 2: We need to obtain the high severity work list of the above obtained object.

Action 2: Response[{"tool_name": "works_list", "arguments": [{"argument_name": "ticket.rev_org", "argument_value": "\$\$PREV[0]"}]}]

Observation 2: We have called works_list and obtained tickets of the object (customer) UltimateCustomer.

Thought 3: We need to summarize the obtained work list.

Action 3: Response[{"tool_name": "summarize_objects", "arguments": [{"argument_name": "objects", "argument_value": "\$\$PREV[1]"}]}]

Observation 3: We have called summarize_objects with the result obtained from works_list.

Thought 4: After this, we have summarized the high severity tickets from the customer UltimateCustomer.

Action 4: Finish[" "]

Data Generation

For testing using **ToolEval**, the input for evaluation was given in the following format (2 data points are shown below) with arguments

- **qid:** The id of the data point
- **question:** Query for which an answer is expected
- **gnd:** The ground truth written manually
- **mod:** The output from ReAct framework

```
[{"qid": 1, "question": "Summarize work items similar to don:core:dvr-v-us-1:devo/0:issue/1 and what is the meaning of life?", "gnd": "Thought 1: This is a question related to summarizing work items similar to don:core:dvr-v-us-1:devo/0:issue/1 and the meaning of life.\nAction 1: Response[\"{\\\"tool_name\\\": \\\"get_similar_work_items\\\", \\\"arguments\\\": [{\\\"argument_name\\\": \\\"work_id\\\", \\\"argument_value\\\": \\\"don:core:dvr-v-us-1:devo/0:issue/1\\\"}]}\\\"]\nObservation 1: done\nThought 2: We need to summarize the obtained work items.\nAction 2: Response[\"{\\\"tool_name\\\": \\\"summarize_objects\\\", \\\"arguments\\\": [{\\\"argument_name\\\": \\\"objects\\\", \\\"argument_value\\\": \\\"$$PREV[0]\\\"}]}\\\"]\nObservation 2: done\nThought 3: After this, we have summarized the work items similar to don:core:dvr-v-us-1:devo/0:issue/1.\nAction 3: Finish[\\\" \\\"]", "mod": "ReAct framework output"}]
```

```

"mod": "Thought 1: This pertains to summarizing tasks akin to don:core:dvr-us-1:devo/0:issue/1 and exploring the significance of existence.\nAction 1: Response[\"{\\\"tool.name\\\": \\\"get.similar.work.items\\\", \\\"arguments\\\": {\\\"argument.name\\\": \\\"work.id\\\", \\\"argument.value\\\": \\\"don:core:dvr-us-1:devo/0:issue/1\\\"}}\\\"]\nObservation 1: done\nThought 2: It is necessary to provide a summary of the obtained work items.\nAction 2: Response[\"{\\\"tool.name\\\": \\\"summarize.objects\\\", \\\"arguments\\\": {\\\"argument.name\\\": \\\"objects\\\", \\\"argument.value\\\": \\\"$PREV[0]\\\"}}\\\"]\nObservation 2: done\nThought 3: Subsequently, we have summarized the work items in a fashion similar to don:core:dvr-us-1:devo/0:issue/1.\nAction 3: Finish[\" \\\"]",
"win": true
}
{
"qid": 2,
"question": "Question: Summarize issues similar to don:core:dvr-us-1:devo/0:issue/1",
"gend": "Thought 1: This is a question related to Summarizing a list of objects.\nAction 1: get.similar.work.items[work.id=don:core:dvr-us-1:devo/0:issue/1]\nObservation 1: Called the function get.similar.work.items and obtained similar work items\nThought 2: We need to summarize the obtained work items.\nAction 2: summarize.objects[objects=$PREV[0]]\nObservation 2: We have called summarize.objects with the result obtained from get.similar.work.items\nThought 3: After this, we have summarized the work items.\nAction 3: Finish[\"{\\\"tool.name\\\": \\\"get.similar.work.items\\\", \\\"arguments\\\": {\\\"argument.name\\\": \\\"work.id\\\", \\\"argument.value\\\": \\\"don:core:dvr-us-1:devo/0:issue/1\\\"}}\\\", {\\\"tool.name\\\": \\\"summarize.objects\\\", \\\"arguments\\\": {\\\"argument.name\\\": \\\"objects\\\", \\\"argument.value\\\": \\\"$PREV[0]\\\"}}\\\"]\n",
"mod": "Thought 1: This is a question related to summarizing high severity tickets from the customer UltimateCustomer.\nAction 1: Response[\"{\\\"tool.name\\\": \\\"search.object.by.name\\\", \\\"arguments\\\": {\\\"argument.name\\\": \\\"query\\\", \\\"argument.value\\\": \\\"UltimateCustomer\\\"}}\\\"]\nObservation 1: Called the function search.object.by.name and obtained the object of UltimateCustomer.\nThought 2: We need to obtain the high severity work list of the above obtained object.\nAction 2: Response[\"{\\\"tool.name\\\": \\\"works.list\\\", \\\"arguments\\\": {\\\"argument.name\\\": \\\"ticket.rev.org\\\", \\\"argument.value\\\": \\\"$PREV[0]\\\"}}\\\"]\nObservation 2: We have called works.list and obtained tickets of the object (customer) UltimateCustomer.\nThought 3: We need to summarize the obtained work list.\nAction 3: Response[\"{\\\"tool.name\\\": \\\"summarize.objects\\\", \\\"arguments\\\": {\\\"argument.name\\\": \\\"objects\\\", \\\"argument.value\\\": \\\"$PREV[1]\\\"}}\\\"]\nObservation 3: We have called summarize.objects with the result obtained from works.list.\nThought 4: After this, we have summarized the high severity tickets from the customer UltimateCustomer.\nAction 4: Finish[\" \\\"]",
"win": false
}
}

```

The results and inferences on data generated using above process is mentioned in Section 3

2.3 ToolLLM with ChatGPT backbone

2.3.1 API-Retrieval

We borrowed ToolLLMs [4] API Retrieval where we first narrow down our search space from all the given pool of APIs. We select top-k APIs which are relevant to the given query and use them to answer the given question. We use the Sentence Transformer Model, which maps sentences and paragraphs to a 768 dimensional dense vector space. We then use cosine similarity to fetch the top-k embeddings of the API's relevant to the given query. We made sure that for a given API, the sentence input to the Sentence Transformer consists of the API name, description, required parameters, optional parameters, and parameters description. With this we make sure scaling our API dataset won't be an issue even in the later stages.

```

process[0] using task 0/17 task_id=1
Retrieving...
[{'corpus_id': 1, 'score': 0.3104080855846405}, {'corpus_id': 7, 'score': 0.06350993365049362}, {'corpus_id': 5, 'score': 0.044816844165325165}, {'corpus_id': 2, 'score': 0.040305137634277344}, {'corpus_id': 4, 'score': -0.0032262448221445084}, {'corpus_id': 8, 'score': -0.00540948286652565}, {'corpus_id': 3, 'score': -0.00562678650021553}, {'corpus_id': 0, 'score': -0.030704565346240997}, {'corpus_id': 6, 'score': -0.11529131978750229}]

```

Figure 3: API-Retrieval's output which is sent to the later modules

2.3.2 Depth First Search-based Decision Tree (DFSDT)

ToolLLM highlighted that previous works didn't fully evicit the capabilities of LLMs by adopting either CoT or React for model Reasoning hence, failing to handle complex tasks. To this end, they proposed Depth First Search-based Decision Tree (DFSDT) solving the main issues of ReAct or CoT which are error propagation and limited exploration. In DFSDT, we construct a decision tree to expand the search space, increasing the possibility of finding a valid path. For searching through the Decision tree, we used Depth First Search (DFS).

In our implementation, where we use ChatGPT as backbone, we exploit the function-calling feature provided by the OpenAI's API to accessing Chat-GPT. That lets up pass the list of functions to Chat-GPT, making our function calls correct in most of the cases.

```
[
  {
    "tool_name": "get_similar_work_items",
    "arguments": [
      {
        "argument_name": "work_id",
        "argument_value": "WK-001"
      }
    ]
  },
  {
    "tool_name": "summarize_objects",
    "arguments": [
      {
        "argument_name": "objects",
        "argument_value": "$$PREV[0]"
      }
    ]
  }
]

[
  {
    "tool_name": "get_similar_work_items",
    "arguments": [
      {
        "argument_name": "work_id",
        "argument_value": "WK-001"
      }
    ]
  },
  {
    "tool_name": "summarize_work_items",
    "arguments": [
      {
        "argument_name": "objects",
        "argument_value": "$$PREV1"
      }
    ]
  }
]
```

Figure 4: Query: Summarize issues similar to WK-001. [Observed output (right) same as expected output (left)]

```
[
  {
    "tool_name": "search_object_by_name",
    "arguments": [
      {
        "argument_name": "query",
        "argument_value": "Cust123"
      }
    ]
  },
  {
    "tool_name": "works_list",
    "arguments": [
      {
        "argument_name": "ticket.rev_org",
        "argument_value": "$$PREV[0]"
      },
      {
        "argument_name": "ticket.severity",
        "argument_value": "high"
      },
      {
        "argument_name": "ticket.source_channel",
        "argument_value": "slack"
      }
    ]
  },
  {
    "tool_name": "summarize_objects",
    "arguments": [
      {
        "argument_name": "objects",
        "argument_value": "$$PREV[1]"
      }
    ]
  }
]

[
  {
    "tool_name": "works_list",
    "arguments": [
      {
        "argument_name": "owned_by",
        "argument_value": "Cust123"
      },
      {
        "argument_name": "ticket_severity",
        "argument_value": "high"
      },
      {
        "argument_name": "ticket_source_channel",
        "argument_value": "slack"
      }
    ]
  },
  {
    "tool_name": "works_list",
    "arguments": [
      {
        "argument_name": "owned_by",
        "argument_value": "$$PREV0"
      },
      {
        "argument_name": "ticket_severity",
        "argument_value": "$$PREV0"
      },
      {
        "argument_name": "ticket_source_channel",
        "argument_value": "$$PREV0"
      }
    ]
  },
  {
    "tool_name": "summarize_work_items",
    "arguments": [
      {
        "argument_name": "objects",
        "argument_value": "$$PREV1"
      }
    ]
  }
]
```

Figure 5: Query: List all high severity tickets coming in from slack from customer Cust123 and generate a summary of them. [Observed output (right) same as expected output (left)]

2.4 Few-shot Prompting

As a continuation of experiments mentioned in the mid-term report, using **ToolQA's GPT-3.5-Turbo model**, we tried the **one-shot prompting** by trying out various prompts with the existing tools and a seed example and obtained answers to all the questions, few of them which are shown below. These were the best results among all the prompts we tried (path of prompt: in the following format (below) ToolQA+GPT/new_prompt.txt).

<pre>"answer": [{ "tool_name": "create_actionable_tasks_from_text", "arguments": [{ "argument_name": "text", "argument_value": "T" }] }, { "tool_name": "get_sprint_id", "arguments": [] }, { "tool_name": "add_work_items_to_sprint", "arguments": [{ "argument_name": "work_ids", "argument_value": "\$\$PREV[0]" }, { "argument_name": "sprint_id", "argument_value": "\$\$PREV[1]" }] }], }</pre>	<pre>"answer": [{ "tool_name": "create_actionable_tasks_from_text", "arguments": [{ "argument_name": "text", "argument_value": "T" }] }, { "tool_name": "get_sprint_id", "arguments": [] }, { "tool_name": "add_work_items_to_sprint", "arguments": [{ "argument_name": "work_ids", "argument_value": "\$\$PREV[0]" }, { "argument_name": "sprint_id", "argument_value": "\$\$PREV[1]" }] }], }</pre>
---	---

Figure 6: **Correct results:** Observed output (right) same as expected output (left)

<pre>"answer": [{ "tool_name": "get_similar_work_items", "arguments": [{ "argument_name": "work_id", "argument_value": "TKT-123" }] }, { "tool_name": "create_actionable_tasks_from_text", "arguments": [{ "argument_name": "text", "argument_value": "\$\$PREV[0]" }] }, { "tool_name": "prioritize_objects", "arguments": [{ "argument_name": "objects", "argument_value": "\$\$PREV[1]" }] }], }</pre>	<pre>"answer": [{ "tool_name": "get_similar_work_items", "arguments": [{ "argument_name": "work_id", "argument_value": "TKT-123" }] }, { "tool_name": "summarize_objects", "arguments": [{ "argument_name": "objects", "argument_value": "\$\$PREV[0]" }] }, { "tool_name": "create_actionable_tasks_from_text", "arguments": [{ "argument_name": "text", "argument_value": "\$\$PREV[1].summary" }] }, { "tool_name": "prioritize_objects", "arguments": [{ "argument_name": "objects", "argument_value": "\$\$PREV[2]" }] }], }</pre>
---	---

Figure 7: **Incorrect results:** Observed output (right) different from expected output (left)

For correctness of the output, if the order of tools obtained by the model is same as the order of tools in the ground truth, it is considered correct output. The argument values are not considered for the correctness. The

reported accuracy is defined as

$$\frac{\text{Correct answers}}{\text{Total number of data points}} \quad (1)$$

Results: Accuracy is **62.5%**.

Challenges: This approach does not work well in cases of compound questions, multiple questions and comparisons with multi-reasoning approach, since the model is just trying to recognise some patterns based on the provided prompt, tools' descriptions and seed examples. The argument values obtained in some cases differ slightly from the ground truth argument values due to which accuracy decreases.

3 Results and Inferences

3.1 ToolEval

To assess the capabilities of the models, the inference is based on the outputs and the process by which the solution is obtained. Following the guidelines of the evaluator ToolEval, backed up by ChatGPT, ToolEval achieves a high correlation with human evaluation and provides a robust, scalable, and reliable assessment for machine tool use. It comprises two key metrics:

- **Pass Rate:** Measures the language model's ability to successfully execute an instruction.
- **Win Rate:** Compares the quality and usefulness of two solution paths.

With ToolEval as motivation, our inference is based on accuracy and path preference. Accuracy is calculated by comparing the expected solution and the model output. Path preference is the path preferred by the language model when given two paths to achieve the goal, which comprises the human-annotated path to arrive at a solution and the path the model outputs to arrive at the solution.

For instance, for the query:

"Get all work items similar to TKT-123, summarize them, create issues from that summary, and prioritize them"
the human-annotated path in ToolLLM evaluation will be:

```
get_similar_work_items → create_actionable_tasks_from_text → prioritize_objects
```

And for the ReAct inference, a sequence of thought, action, and observation is fed. The inference is made on the outputs of ReAct and ToolLLM and are presented below. The inference was made on 8 questions on ToolLLM and 5 questions with ReAct. Each data comprises of query, human annotated path and the model output path along with the correctness of the output.

Accuracy is calculated considering the correctness with respect to the relevant tools picked by the model.

Preference is obtained by providing an instruction and two solution paths to ChatGPT evaluator and obtain its preference (i.e., which one is better). We pre-define a set of criteria for both metrics and these criteria are organized as prompts for our ChatGPT evaluator. We evaluate multiple times based on ChatGPT to improve the reliability and then take the average.

- **Path 1** represents the path suggested by the model.
- **Path 2** represents the path annotated by human.

Question	Preferred Path	Path 1 (%)	Path 2 (%)	Total Responses
1	1	17 (85.00%)	3 (15.00%)	20
2	2	8 (40.00%)	12 (60.00%)	20
3	2	8 (40.00%)	12 (60.00%)	20
4	2	0 (0.00%)	20 (100.00%)	20
5	1	11 (55.00%)	9 (45.00%)	20
6	1	13 (65.00%)	7 (35.00%)	20
7	2	8 (40.00%)	12 (60.00%)	20
8	2	0 (0.00%)	20 (100.00%)	20
Preferred Counts		65	95	160
Avg Preference		40.62%	59.38%	

Table 1: Inference obtained for ToolLLM outputs

Question	Preferred Path	Path 1 (%)	Path 2 (%)	Total Responses
1	1	17 (85.00%)	3 (15.00%)	20
2	1	14 (70.00%)	6 (30.00%)	20
3	1	14 (70.00%)	6 (30.00%)	20
4	1	15 (75.00%)	5 (25.00%)	20
5	1	15 (75.00%)	5 (25.00%)	20
6	1	17 (85.00%)	3 (15.00%)	20
7	1	19 (95.00%)	1 (5.00%)	20
8	1	17 (85.00%)	3 (15.00%)	20
Preferred Counts		128	32	160
Avg Preference		100.00%	0.00%	

Table 2: Inference obtained for CoT outputs

Question	Preferred Path	Path 1 (%)	Path 2 (%)	Total Responses
1	1	22 (73.33%)	8 (26.67%)	30
2	1	18 (60.00%)	12 (40.00%)	30
3	1	17 (56.67%)	13 (43.33%)	30
4	1	23 (76.67%)	7 (23.33%)	30
5	2	14 (46.67%)	16 (53.33%)	30
Preferred Counts		94	56	150
Avg Preference		62.67%	37.33%	

Table 3: Inference obtained for ReAct outputs

For the calculation of Correctness in Tool selection, the models are fed with 8 questions provided in the Problem statement, and the observation is made as follows:

Tool	Correctness in Tool selection	Preferred Path
ToolLLM	37.5%	Human Annotated
CoT	87.5%	Model path
ReAct	75%	Model path

Table 4: Model Accuracy and Preferred Paths

From the above results we see the suggested path is chosen in two experiments in the evalutaion suggesting the correctness of the model. The evalutaion prompt is also instructed to prefer the model path if model path and the human annotated path are the same. Hence the model path is preferred in case of ReAct and CoT whereas

TooLLM solution pathway is not preferred. This is also evident with the accuracy w.r.t given data which is higher for ReAct and CoT.

3.1.1 Difference between CoT and ReAct

Chain-of-thought (CoT) prompting has shown the capabilities of LLMs to carry out reasoning traces to generate answers to questions involving arithmetic and commonsense reasoning, among other tasks. But its lack of access to the external world or inability to update its knowledge can lead to issues like fact hallucination and error propagation. In CoT prompting, the entire chain of actions or instructions is generated at once. This means that the system receives a complete set of instructions or thoughts in advance, and it processes them sequentially without considering external observations or responses during the execution.

ReAct is a general paradigm that combines reasoning and acting with LLMs. ReAct prompts LLMs to generate verbal reasoning traces and actions for a task. This allows the system to perform dynamic reasoning to create, maintain, and adjust plans for acting while also enabling interaction to external environments to incorporate additional information into the reasoning. In ReAct prompting, each action is decided based on the previous thought, action, and real-time observations or responses received from the environment or API. It involves an iterative process where the system responds to the environment's feedback, observes the outcomes of its previous actions, and then decides the next action accordingly.

3.1.2 Analysis of Model Usage: Tokens and Time

The analysis on token usage and time taken to give right answer is made on testing questions given in the problem statement and the average is reported here. The number of input tokens vary for the ReAct and TooLLM model since the intermediate steps depend on the context. However in the case of Chain of thought prompting input prompt is the same except for the question to be asked, so the size of input prompt is 1773 tokens on an average.

Model	Time Taken (s)
TooLLM	33.11
CoT	5.67
ReAct	21.51
ToolQA + GPT	11.15

Table 5: Time Analysis

This analysis shows that CoT returns results in least time.

3.2 Reasoning on the Issues with TooLLM (ChatGPT Backbone)

Despite showing promising results for real world api's, we had a hard time in making the framework generate correct chain of api calls for the following reasons:

- **Problem with question asking to Summarize work items:** With the usage of powerful Large Language Models (LLMs) such as GPT-4, the LLM is using its knowledge of the work items (their ID or description) and the knowledge of summarize_work_items api (based on description) to summarize the items instead of calling the api summarize_work_items. Even though it was able to summarize the work items perfectly using existing knowledge, it was breaking our chain of api calls thus, leading to the wrong output.
- **Inconsistency in the return type:** The TooLLM framework was trained on 16000+ Real world RestAPI's. They were trained on returning outputs in json format. But for our use case, we needed the outputs as a list since, the function calls take list of objects as input. Since TooLLM takes decisions dynamically based on the response from LLM, the output format in which the response was delivered was inconsistent, making it hard for us in taking the arguments for the function calls.

- **Inability to revert to an old state where a API makes state changes** We are interested in adding work items with some requirements to the current sprint. ToolLLM will retrieve the work items with the given requirements. Let's assume that while giving requirements, it made some errors. So we got the wrong work items. We now add the wrong work items to the sprint. But if ToolLLM now realizes that the obtained work items are wrong, it will go back and start getting the work items again and add them to current sprint again. It doesn't remove the already wrongly added work items.

4 New Tools (Bonus Part)

For the bonus part, following new tools were created in order to handle queries which require combining outputs of multiple functions using iterations, conditional logic, mathematical operations. These has been implemented in ReAct.

Tool	Description	Argument Name	Argument Description	Argument Type
select_first_k_tasks	Returns the first k objects from a list of objects	objects	List of objects from which first k are to be selected	array of objects
select_last_k_tasks	Returns the last k objects from a list of objects	objects	List of objects from which last k are to be selected	array of objects
compare_priorities	Compares the priorities of two objects and returns the higher priority object	object1, object2	Two lists of objects whose priorities are to be compared	array of objects
count_tasks	Counts and returns the number of tasks in a list of objects	objects	List of objects which are to be counted	array of objects
average_priority	Calculates and returns the average priority of a list of objects	arg1	List of objects whose average priority is to be computed	array of objects

Table 6: New Tools

Below are a few queries which illustrate the use of above created tools.

Query	Output
<p>Prioritize my P0 issues and add the first 5 to the current sprint</p>	<pre>[{ "tool.name": "whoami", "arguments": [] }, { "tool.name": "works.list", "arguments": [{ "argument.name": "issue.priority", "argument.value": "p0" }, { "argument.name": "owned.by", "argument.value": "\$\$PREV[0]" }] }, { "tool.name": "prioritize.objects", "arguments": [{ "argument.name": "objects", "argument.value": "\$\$PREV[1]" }] }, { "tool.name": "select.first.k.tasks", "arguments": [{ "argument.name": "objects", "argument.value": "\$\$PREV[2]" }] }, { "tool.name": "get.sprint.id", "arguments": [] }, { "tool.name": "add.work.items.to.sprint", "arguments": [{ "argument.name": "work.ids", "argument.value": "\$\$PREV[3]" }, { "argument.name": "sprint.id", "argument.value": "\$\$PREV[4]" }] }]</pre>
<p>Get all work items owned by me, summarize them, and calculate the average priority</p>	<pre>[{ "tool.name": "who.am.i", "arguments": [] }, { "tool.name": "works.list", "arguments": [{ "argument.name": "owned.by", "argument.value": "\$\$PREV[0]" }] }, { "tool.name": "summarize.objects", "arguments": [{ "argument.name": "objects", "argument.value": "\$\$PREV[1]" }] }, { "tool.name": "average.priority", "arguments": [{ "argument.name": "objects", "argument.value": "\$\$PREV[2]" }] }]</pre>

<p>Get my highest priority issue between TKT-123 and TKT-456</p>	<pre>[{ "tool_name": "who.am.i", "arguments": [] }, { "tool_name": "get.similar.work.items", "arguments": [{ "argument_name": "work.id", "argument_value": "TKT-123" }] }, { "tool_name": "get.similar.work.items", "arguments": [{ "argument_name": "work.id", "argument_value": "TKT-456" }] }, { "tool_name": "compare.priorities", "arguments": [{ "argument_name": "object1", "argument_value": "\$\$PREV[1][0]" }, { "argument_name": "object2", "argument_value": "\$\$PREV[2][0]" }] }]</pre>
<p>Count the number of tasks in the current sprint</p>	<pre>[{ "tool_name": "get.sprint.id", "arguments": [] }, { "tool_name": "works.list", "arguments": [{ "argument_name": "sprint.id", "argument_value": "\$\$PREV[0]" }] }, { "tool_name": "count.tasks", "arguments": [{ "argument_name": "objects", "argument_value": "\$\$PREV[1]" }] }]</pre>

5 Future Work

JuDeC [5], a novel approach enhancing large language models (LLMs) with self-judgment capabilities for autonomous decision-making using an Elo-based Self-Judgment Mechanism can be implemented from scratch given more time as this approach not only provides higher-quality solutions but also demonstrates cost efficiency, making it an effective tool for LLMs in multi-step decision-making processes. In addition to Chain of Thought prompting different prompting techniques like Tree of thoughts and Graph of Thoughts can be employed.

6 Conclusion

This report presents our implementation results and experiments based on the literature survey we had done in the mid-term report. Among the GPT, CoT, DFSDT and ReAct frameworks, **CoT gave the best accuracy of 87.5%**. The superior performance of CoT over GPT, DFSDT, and ReAct frameworks may be attributed to its advanced

context-based understanding, effective handling of diverse input data, and optimized learning mechanisms. Its ability to leverage contextual information, adapt to varying scenarios, and employ robust training strategies likely contributes to its higher accuracy.

References

- [1] Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models, 2023.
- [2] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. React: Synergizing reasoning and acting in language models. *arXiv preprint arXiv:2210.03629*, 2022.
- [3] Yuchen Zhuang, Yue Yu, Kuan Wang, Haotian Sun, and Chao Zhang. Toolqa: A dataset for llm question answering with external tools, 2023.
- [4] Yujia Qin, Shihao Liang, Yining Ye, Kunlun Zhu, Lan Yan, Yaxi Lu, Yankai Lin, Xin Cong, Xiangru Tang, Bill Qian, Sihan Zhao, Lauren Hong, Runchu Tian, Ruobing Xie, Jie Zhou, Mark Gerstein, Dahai Li, Zhiyuan Liu, and Maosong Sun. Toolllm: Facilitating large language models to master 16000+ real-world apis, 2023.
- [5] Yining Ye, Xin Cong, Yujia Qin, Yankai Lin, Zhiyuan Liu, and Maosong Sun. Large language model as autonomous decision maker. *arXiv preprint arXiv:2308.12519*, 2023.